

ANKIT SAXENA

Senior Data Engineer

Pune, India | +91 7500335301 | ankit282k@gmail.com | linkedin.com/in/ankit282k | github.com/ankit282k

PROFESSIONAL SUMMARY

Data Engineer with 2+ years of experience architecting and productionizing large-scale ETL/ELT pipelines, ML-ready data platforms, and GenAI-enabled microservices across PySpark, Databricks, Snowflake, and Delta Lake. AWS Certified Cloud Practitioner with hands-on experience delivering cost- and performance-optimized data infrastructure processing 30M+ records daily, cutting pipeline latency by up to 45% and infrastructure cost by 22%. Track record of owning projects end-to-end — from data modeling and pipeline design through CI/CD deployment and production monitoring — across retail, enterprise AI, and franchise-management domains. Seeking a Senior Data Engineer role to drive scalable data architecture, streaming/batch pipeline strategy, and MLOps best practices.

CORE COMPETENCIES

Data Engineering: PySpark, Apache Spark, Databricks, Delta Lake, Data Lake, Data Warehousing, Snowflake, SQL, Big Data, ETL/ELT Pipelines, Data Modeling, Medallion Architecture (Bronze/Silver/Gold), Star Schema, Kafka (Streaming), Hive, Batch & Streaming Processing, Schema Evolution, Partition Tuning, Data Governance, Data Quality

Cloud & DevOps: AWS (S3, EC2, IAM, CloudWatch, Redshift) — Certified Cloud Practitioner, Docker, GitHub Actions, CI/CD, Linux, Version Control (Git)

Programming & Orchestration: Python, SQL, FastAPI, REST APIs, Airflow, Prefect, Git, Agile/Scrum

ML / GenAI Enablement: MLflow, Databricks AutoML, AutoML, OpenAI SDK, LLM Integration, Time Series Forecasting (XGBoost, ARIMA), Feature Engineering, HuggingFace, Prompt Engineering, LLM Microservices, MLOps

PROFESSIONAL EXPERIENCE

AI/ML Engineer I | Simplifx Systems (Client: Vijay Sales)

Jan 2026 – Present

- Architect and maintain end-to-end Databricks pipelines using PySpark, AWS, and MongoDB to process large-scale retail sales and inventory data supporting enterprise analytics and demand forecasting.
- Design and manage a centralized Data Lake on AWS to consolidate multi-source retail sales and inventory data, enabling scalable storage, governance, and self-service analytics across teams.
- Lead implementation of 21-day item–store level time-series demand forecasting to drive inventory planning and stock replenishment, directly improving product availability.
- Leverage Databricks AutoML to automate training, evaluation, and selection of forecasting models, cutting manual experimentation time and accelerating the model development lifecycle.
- Build GenAI-powered analytics assistants and automation workflows using the OpenAI SDK, integrating LLM-based reasoning with internal data pipelines to support natural-language reporting and insights.
- Own experiment tracking, model versioning, and production deployment using MLflow, establishing ongoing monitoring for forecasting models in live workflows.

Associate Software Developer | Bosch Global Software Technologies

Feb 2025 – Jan 2026

- Designed and optimized PySpark-based ETL pipelines for internal AI platforms, reducing processing latency by 28% and storage costs by 22%.
- Built cloud-ready, production-grade data workflows supporting ML and GenAI use cases using Databricks, Delta Lake, and distributed Spark jobs.
- Engineered LLM-powered microservices with FastAPI and HuggingFace, integrating secure data pipelines that improved team productivity by 32%.
- Implemented CI/CD pipelines with GitHub Actions and Docker, automating testing and deployment and reducing production issues by 80%.

Data Engineer Trainee | FranConnect Pvt. Ltd.

Sep 2024 – Feb 2025

- Built and maintained large-scale ETL pipelines using PySpark, Hive, and SQL, processing 30M+ records daily with a 45% performance improvement.

- Implemented incremental ingestion, schema evolution, partition tuning, and caching strategies, reducing Spark job runtime by 22%.
- Automated data pipeline deployments using Docker and GitHub Actions, cutting integration failures by 80%.

Data Scientist Trainee | Dyno India Pvt. Ltd.

Feb 2023 – Aug 2023

- Built deep learning pipelines using CNN and LSTM architectures, improving validation metrics by 30%.
- Implemented recommender systems using Collaborative Filtering and FISM, improving recommendation precision by 25%.
- Executed complete ML workflows from EDA and feature engineering through model deployment and documentation.

KEY PROJECTS

End-to-End Databricks ETL Pipeline (PySpark, Delta Lake, AWS)

- Designed a Bronze → Silver → Gold Medallion Architecture on Databricks and Delta Lake with incremental loads, validation, and transformation jobs for large datasets.
- Used AWS S3 as the storage layer with IAM-based access control; improved pipeline runtime performance by 42%.

Snowflake Automated Data Warehouse (Snowpipe, Streams, Tasks)

- Built automated ingestion pipelines using Snowpipe from cloud object storage and designed fact/dimension tables with SCD logic for analytics use cases.
- Implemented Streams and Tasks for continuous incremental data processing.

Time Series Forecasting with MLflow

- Developed forecasting pipelines using XGBoost and ARIMA, achieving a 21% MAPE improvement, with full experiment and model tracking via MLflow.

CERTIFICATIONS

- AWS Certified Cloud Practitioner (CLF-C02)

EDUCATION

Bachelor of Technology – Electronics & Communication Engineering

Chandigarh Engineering College | CGPA: 7.93 | 2023